

# APM 598: Homework 3 (03/26)

## 1 n-gram models

### Ex 1.

- a) Load and tokenize the text attached `'Plato_Republic.txt'`.  
Put all the words in lower case to regroup words like 'The' and 'the'.  
Compute the total number of words  $N$  in the text and the number of unique words (size of the vocabulary).

- b) Build a uni-gram. Deduce the 5 most common words with **at least 8 characters**.  
*Hint: use the method `'most_common'` on an object `'nltk.FreqDist'`.*

- c) Build a bi-gram and define a function that given two words  $(x_1, x_2)$  compute the probability:

$$\mathbb{P}(x_2|x_1) = \frac{\#\{(x_1, x_2)\}}{\#\{x_1\}}$$

where  $\#$  denotes the number of occurrences of the word (or pair of words) in the corpus.

- d) Deduce the so-called perplexity of the bi-gram model defined as:

$$PP = \left( \prod_{k=1..(N-1)} \mathbb{P}(x_{k+1}|x_k) \right)^{-\frac{1}{N-1}}$$

where  $N$  denotes the total number of words in the corpus.

## 2 Recurrent Neural Networks

### Ex 2.

The goal of this exercise is to experiment with a simple Recurrent Neural Network (RNN) model for predicting letters. We only consider four letters "h", "e", "l" and "o" that we embed in  $\mathbb{R}^4$ :

$$\text{"h"} \rightarrow \begin{pmatrix} 1 \\ 0 \\ 0 \\ 0 \end{pmatrix}, \quad \text{"e"} \rightarrow \begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \end{pmatrix}, \quad \text{"l"} \rightarrow \begin{pmatrix} 0 \\ 0 \\ 1 \\ 0 \end{pmatrix}, \quad \text{"o"} \rightarrow \begin{pmatrix} 0 \\ 0 \\ 0 \\ 1 \end{pmatrix}.$$

We consider a RNN with hidden states  $\mathbf{h}_t$  in  $\mathbb{R}^2$ :

$$\begin{cases} \mathbf{h}_t &= \tanh(R\mathbf{h}_{t-1} + A\mathbf{x}_t) \\ \mathbf{y}_t &= B\mathbf{h}_t \end{cases} \quad (1)$$

where  $A \in \mathcal{M}_{2,4}(\mathbb{R})$ ,  $R \in \mathcal{M}_{2,2}(\mathbb{R})$  and  $B \in \mathcal{M}_{4,2}(\mathbb{R})$  (e.g.  $A$  is a  $2 \times 4$  matrix).

- a) Given the input "hello" (i.e.  $\mathbf{x}_1 = (1, 0, 0, 0), \dots, \mathbf{x}_5 = (0, 0, 0, 1)$ ), the initial state  $\mathbf{h}_0 = (0, 0)$  and the matrices:

$$A = \begin{bmatrix} 1 & -1 & -1/2 & 1/2 \\ 1 & 1 & -1/2 & -1 \end{bmatrix}, \quad R = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad B = \begin{bmatrix} 1 & 1 \\ 1/2 & 1 \\ -1 & 0 \\ 0 & -1/2 \end{bmatrix},$$

find the output  $\mathbf{y}_1, \dots, \mathbf{y}_5$  and deduce the predicted characters (see figure 1).

- b) Find matrices  $A$ ,  $R$ ,  $B$  such that the predicted characters are "olleh".




Figure 1: Predictions of a vanilla RNN. After encoding the letters (e.g. "h") into vectors (e.g.  $\mathbf{x}_1 = (1, 0, 0, 0)$ ), the network performs the operations described in eq. (1) to estimate a vector prediction (e.g.  $\mathbf{y}_1$ ). The 'letter' predicted is chosen as the index of the output with the largest value (i.e. find the hot vector the closest to (softmax) of  $\mathbf{y}_1$ ).

**Ex 3.** [vanishing/exploding gradient]

We would like to illustrate one of the issue with *vanilla RNN*, namely the vanishing or exploding gradient phenomenon. Rather than computing the gradient of the loss function, we simply are going to investigate how a small perturbation in the input  $\mathbf{x}_1$  will affect the output  $\mathbf{y}_t$  (see figure 2).




Figure 2: To study how a perturbation of  $\mathbf{x}_1$  affects  $\mathbf{y}_t$ , we suppose in this exercise that  $\mathbf{x}_2 = \dots \mathbf{x}_t = \mathbf{0}$  and  $\mathbf{h}_0 = \mathbf{0}$ . Due to the iterations of the matrix  $R$  in the estimation of  $\mathbf{y}_t$ , the perturbation of  $\mathbf{x}_1$  could have small or large influence on  $\mathbf{y}_t$ .

We consider a standard RNN defined with three matrices  $A, R, B$  and  $\sigma(x) = \tanh(x)$  (see figure 2).

- a) Compute the differential  $D_{\mathbf{h}_{t-1}} \mathbf{h}_t$ , i.e. compute the differential of the function  $\mathbf{h} \rightarrow \sigma(R\mathbf{h} + A\mathbf{x}_t)$ .  
Deduce that:

$$\|D_{\mathbf{x}_1} \mathbf{y}_t\| \leq \|B\| \cdot \left( \prod_{k=1}^t |\sigma'(R\mathbf{h}_{k-1} + A\mathbf{x}_k)|_\infty \right) \cdot \|R\|^{t-1} \cdot \|A\|, \quad (2)$$

where  $\|\cdot\|$  denotes (any) matrix norm and  $|\sigma'(\mathbf{h})|_\infty = \max(|\sigma'(h_1)|, \dots, |\sigma'(h_d)|)$  where  $d$  is the dimension of the vector  $\mathbf{h}$ .

- b) From now on, we take  $t = 30$  and suppose  $\mathbf{x}, \mathbf{y}, \mathbf{h} \in \mathbb{R}^2$  with:

$$A = B = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, R = \begin{bmatrix} \frac{1}{2} & -1 \\ -1 & \frac{1}{2} \end{bmatrix}, \quad \mathbf{x}_2 = \mathbf{x}_3 = \dots = \mathbf{x}_{30} = \mathbf{h}_0 = \begin{pmatrix} 0 \\ 0 \end{pmatrix}.$$

Denote  $\mathbf{x}_1 = (0, 0)$  and  $\mathbf{y}_{30}$  the output after  $t = 30$  iterations.

Similarly, denote the perturbation  $\mathbf{x}_1^\varepsilon = (\varepsilon, -\varepsilon)$  and  $\mathbf{y}_{30}^\varepsilon$  the output after  $t = 30$  iterations starting from  $\mathbf{x}_1^\varepsilon$ .

Compute and plot (in log-log scale) the difference  $\|\mathbf{y}_{30} - \mathbf{y}_{30}^\varepsilon\|$  for  $\varepsilon \in (10^{-4}, \dots, 10^{-9})$ . Relate the result with eq. (2).

- c) Proceed similarly as b) using  $\mathbf{x}_1 = (2, 1)$  and  $\mathbf{x}_1^\varepsilon = (2 + \varepsilon, 1 - \varepsilon)$ .  
Why does the perturbation have a small effect in this case compare to b)?

*Extra)* Proceed similarly as b) using  $\mathbf{x}_1 = (0, 0)$  and  $\mathbf{x}_1^\varepsilon = (\varepsilon, \varepsilon)$ . Why is the perturbation having a small effect? In general, from a random perturbation, do you expect a small or large effect when  $\mathbf{x}_1 = (0, 0)$ ?